

[High Level Design] Using only Customer Container Registry

- 1. PDD
- 2. Dependency
- 3. Prerequisite
- 4. Caveat
- 5. High Level Design
 - 5.1. Workflow
 - 5.2. Key Configuration Parameters
 - 5.2.1. Configuration Example In Job
 - 5.2.2. Configuration Example in Job Run
 - 5.3. Cluster to Job Mapping
 - 5.4. Dask
 - 5.4.1.1. Building Images
 - 5.4.1.1.1. Launching Cluster
 - 5.4.1.2. Alternate designs (Designs that were considered but not chosen)
 - 5.4.1.2.1. - With a Driver image
 - 5.4.1.2.2. Starting workers directly from ADS optcl
 - 5.4.1.2.3. Provide Launch scripts.
 - 5.4.1.2.4. Scheduler starting n workers.

1. PDD

[ADS Distributed Operators PDD](#)

2. Dependency

- Customer has environment to build docker images
- Customer has access to publish docker image in oci container registry
- ML Jobs support BYOD
- ML Jobs supports for Mounted file system (Required for Horovod)
- ML Jobs Firewall does not block the required port numbers

3. Prerequisite

- User will create a regional private subnet with access to NAT GW or Service GW
- User will update the security list to open ports as required by the distributed framework - The port list will be published
- User will setup resource principal policy such that
 - ML Jobs can pull the container
 - ML Jobs can create new Jobs on the same subnet and compartment
 - ML Jobs can access/use user provided subnet
 - ML Jobs can access object storage (optional, this is required if user has to download/access data from object storage during training time. This might also be useful if user wants to save the model in the object storage)

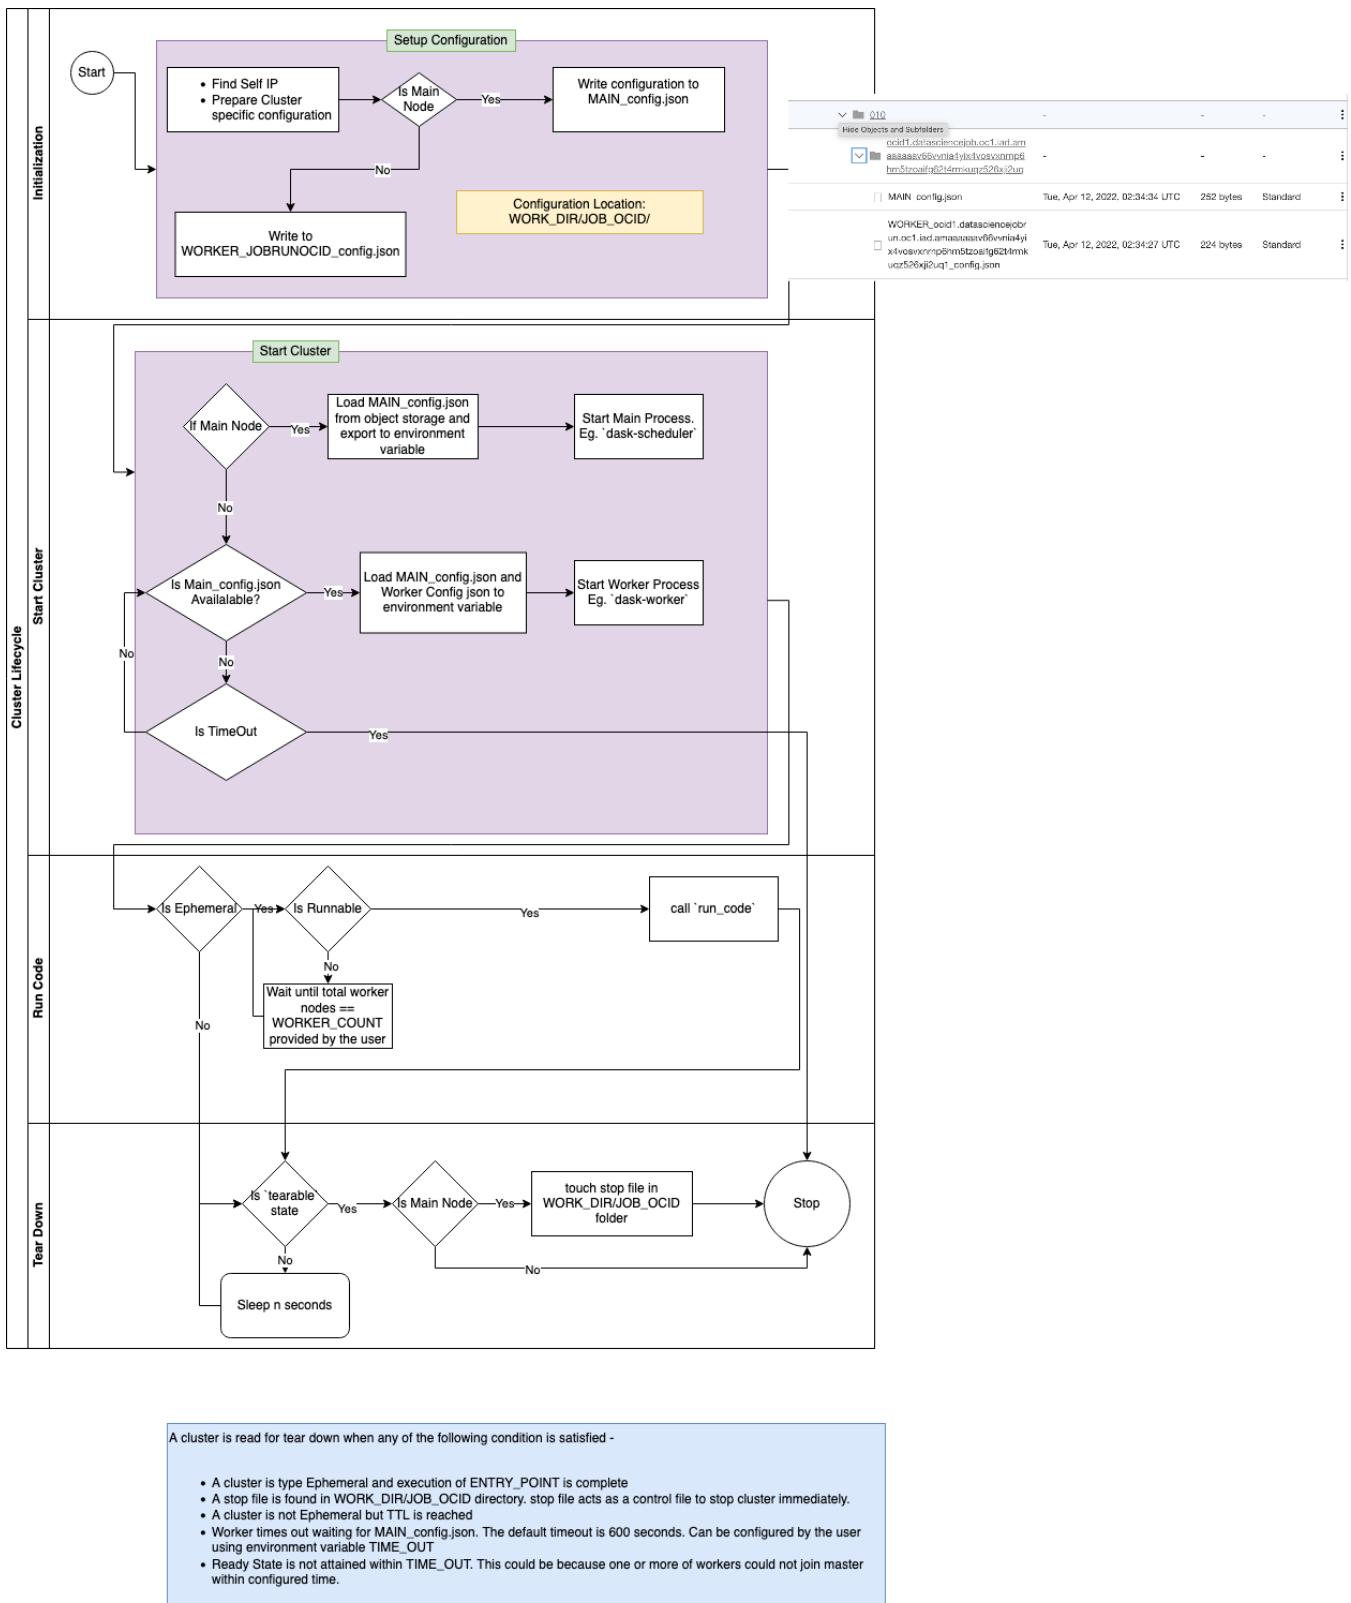
4. Caveat

The main goal of this exercise is to allow users to run distributed training using Jobs. The exact nature of that interaction is less important than the ability of the user to be able to run distributed training jobs. While this design document aims to capture possible UX, we should **not** consider to stick to the design, but allow continuous evolution as we research, experiment and gather feedback. From that sense this document will be continuously updated as we change our approach based on our research.

5. High Level Design

5.1. Workflow

https://bitbucket.oci.oraclecorp.com/projects/ODSC/repos/ads-misc/browse/dsc_distributed?at=distributed



5.2. Key Configuration Parameters

Parameter	DEFAULT	REQUIRED	Description

WORK_DIR	NONE	Y	fsspec supported filesystem - oci objectstorage, s3, gitfs, etc
CLUSTER_TYPE	NONE	Y	The bootstrap process uses this information to delegate the cluster initialization process to the ClusterProvider implementation corresponding to CLUSTER_TYPE
WORKER_COUNT	1	N	Number of expected worker nodes. Can be used to check if the cluster is ready for execution
EPHEMERAL	1	N	If 1, then cluster will run the code provided in `ENTRY_SCRIPT` and stops cluster. If 0, the cluster is expected to be long lived.
ENTRY_POINT	NONE	N	The code to run when launched as Ephemeral Cluster
MODE	NONE	Y	The value here is either MAIN or WORKER. Used to identify the behavior of the Node. If set as MAIN, this node becomes the scheduler. If set as WORKER, behaves as worker.
TIMEOUT	600 Seconds	N	TIME_OUT for the cluster to get into ready state or Wait time for Main_config.json

5.2.1. Configuration Example In Job

Job information Default configuration Tags

Command line arguments: -

Maximum runtime (in minutes): -

Custom environment variables:

- nprocs:** 4 [Show](#) [Copy](#)
- nthread:** 1 [Show](#) [Copy](#)
- death_timeout:** 10 [Show](#) [Copy](#)
- WORK_DIR:** oci://mayoor-dev@ociodscdev/daskcluster-testing/001 [Hide](#) [Copy](#)
- WORKER_COUNT:** 1 [Show](#) [Copy](#)
- CLUSTER_TYPE:** DASK [Show](#) [Copy](#)
- EPHEMERAL:** 1 [Show](#) [Copy](#)
- HELLO:** WORLD [Show](#) [Copy](#)
- DOCKER_CUSTOM_IMAGE:** iad.ocir.io/ociodscdev/dask-cluster-testing:dev [Hide](#) [Copy](#)

5.2.2. Configuration Example in Job Run

Job run Runtime configuration Tags

<h3>Default configuration</h3> <p>Command line arguments: -</p> <p>Maximum runtime (in minutes): -</p> <p>Custom environment variables:</p> <ul style="list-style-type: none"> nprocs: 4 Show Copy nthread: 1 Show Copy death_timeout: 10 Show Copy WORK_DIR: oci://mayoor-dev@ociodscdev/daskcluster-te sting/001 Hide Copy WORKER_COUNT: 1 Show Copy CLUSTER_TYPE: DASK Show Copy EPHEMERAL: 1 Show Copy HELLO: WORLD Show Copy DOCKER_CUSTOM_IMAGE: iad.ocir.io/ociodscdev/dask- cluster-testing:dev Hide Copy 	<h3>Override configuration</h3> <p>Command line arguments: -</p> <p>Maximum runtime (in minutes): -</p> <p>Custom environment variables:</p> <ul style="list-style-type: none"> MODE: MAIN Show Copy ENTRY_POINT: /code/printhello.py Show Copy ARGS: 50 Show Copy NAMED_ARGS: 1 Show Copy
--	---

Job run Runtime configuration Tags

<h3>Default configuration</h3> <p>Command line arguments: -</p> <p>Maximum runtime (in minutes): -</p> <p>Custom environment variables:</p> <ul style="list-style-type: none"> nprocs: 4 Show Copy nthread: 1 Show Copy death_timeout: 10 Show Copy WORK_DIR: oci://mayoor-dev@ociodscdev/daskcluster-te sting/001 Hide Copy WORKER_COUNT: 1 Show Copy CLUSTER_TYPE: DASK Show Copy EPHEMERAL: 1 Show Copy HELLO: WORLD Show Copy DOCKER_CUSTOM_IMAGE: iad.ocir.io/ociodscdev/dask- cluster-testing:dev Hide Copy 	<h3>Override configuration</h3> <p>Command line arguments: -</p> <p>Maximum runtime (in minutes): -</p> <p>Custom environment variables:</p> <ul style="list-style-type: none"> MODE: WORKER Show Copy
--	---

5.3. Cluster to Job Mapping

A cluster is defined by Job and the running instances are Job Runs for that Job Definition. A Job is defined everytime a new cluster is created. Though technically, the job definition can be used to lauch new clusters, the management can get complex as there could be old configurations that are there which can interfere with the new cluster. This design proposes that we have a new Job definition for every new cluster launch.

Here is an example of Dask Cluster. Here we have multiple job runs under the same job definition. The worker/scheduler is decided by the MODE = MAIN or WORKER provided in the environment variable.

ACTIVE

<h3>General information</h3> <p>Description: -</p> <p>OCID: ...bsfd2s7j5q Show Copy</p> <p>Created by: mayoora.rao@oracle.com</p> <p>Created on: Sat, Apr 9, 2022, 04:12:37 UTC</p> <p>Job artifact: container.py (258 bytes)</p>	<h3>Infrastructure configuration</h3> <p>Compute instance shape: VM.Standard2.1</p> <p>Storage: 50 GB</p> <p>VCN: DevTest1-iad.vcn</p> <p>Subnet: ads-testing</p>
<h3>Logging configuration</h3> <p>Logging enabled: true</p> <p>Default log group: mayoor-job</p> <p>Enable automatic log creation: false</p>	

Job runs in networks Compartment

Start a job run				
Name	Status	Lifecycle details	Created by	Time accepted
worker	● Succeeded	-	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 04:50:29 UTC
scheduler	● Succeeded	-	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 04:50:15 UTC

Here is view of multiple cluster launches -

Jobs in networks Compartment

Create job				
Name	Status	Created by	Created on	
dask-trial-3	● Active	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 04:12:37 UTC	⋮
dask-trial-2	● Deleted	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 03:21:01 UTC	⋮
dask-trial-2	● Active	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 01:35:18 UTC	⋮
dask-trial-2	● Deleted	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 00:10:16 UTC	⋮
dask-trail-1	● Active	mayoora.rao@oracle.com	Sat, Apr 9, 2022, 00:00:21 UTC	⋮
dask-trail-1	● Deleted	mayoora.rao@oracle.com	Fri, Apr 8, 2022, 22:22:34 UTC	⋮
dask-hello-world	● Active	mayoora.rao@oracle.com	Fri, Apr 8, 2022, 21:41:29 UTC	⋮
dask-hello-world	● Deleted	mayoora.rao@oracle.com	Fri, Apr 8, 2022, 16:02:38 UTC	⋮

WORK_DIR can be reused by the user across the Cluster creation. The runtime configurations will be isolated using the JOB_OCID folder within WORK_DIR provided by the user.

5.4. Dask

5.4.1.1. Building Images

ADS will contain the Docker recipe to help build the docker image required for Starting a dask cluster. These recipes could be either bundled inside ads wheel or it could be published inside [ADS Github](#). The Docker recipe will include -

The Source below is only for illustration purpose so that we get a sense of what pieces will go into the deliverable. The actual content and exact position of the orchestration will evolve as we develop and experiment

- **Dockerfile Source**

Dockerfile

```
FROM ghcr.io/oracle/oraclelinux7-instantclient:19 AS base

RUN rm -rf /var/cache/yum/* && yum clean all && yum install -y gcc mesa-libGL vim iproute net-tools && rm -rf /var/cache/yum/*
RUN curl -L https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh >> miniconda.sh
RUN bash ./miniconda.sh -b -p /miniconda; rm ./miniconda.sh;
ENV PATH="/miniconda/bin:$PATH"

ENV HOME /home/datascience
RUN mkdir -p /etc/datascience
WORKDIR /etc/datascience

COPY env.yaml /opt/env.yaml
RUN conda env create -f /opt/env.yaml --name daskenv && conda clean -afy
ENV PATH="/miniconda/envs/daskenv/bin:$PATH"

RUN /bin/bash -c "source activate daskenv"
RUN conda init bash && source ~/.bashrc && conda activate daskenv
COPY run.py /etc/datascience/run.py
COPY cluster_helper.py /etc/datascience/cluster_helper.py

RUN mkdir -p /etc/datascience/dask
COPY start-master.sh /etc/datascience/dask/start-master.sh
COPY start-worker.sh /etc/datascience/dask/start-worker.sh
COPY run.sh /etc/datascience/run.sh

RUN chmod u+x /etc/datascience/run.sh
RUN chmod u+x /etc/datascience/dask/start-master.sh
RUN chmod u+x /etc/datascience/dask/start-worker.sh

ENTRYPOINT [ "/etc/datascience/run.sh" ]

ARG CODE_DIR
COPY ${CODE_DIR} /code

EXPOSE 3000-3100
EXPOSE 8700-8800
```

- **Scripts to start the cluster**

start-master.sh

```
#!/bin/bash
set -m -e -o pipefail
ifconfig
ip addr
#export MASTER_IP=$(ifconfig ens5 | grep -Po 't \K[\d.]+ | head -1)
#export MASTER_IP=$(ip addr show ens5 | grep -Po 'inet \K[\d.]+')
#export MASTER_IP=$(hostname -i)

#echo $MASTER_IP
#echo $CONFIG_FILE_PATH
export PYTHONPATH=/code
dask-scheduler --host 0.0.0.0 --port 8786 --pid-file main_pid &
export MAIN_PID=$(<main_pid)
echo "main pid $MAIN_PID"
python /code/$ENTRY_SCRIPT
```

start-worker.sh

```
#!/bin/bash
set -m -e -o pipefail

export PYTHONPATH=/code

echo "num procs"
echo $nprocs
echo "death timeout $death_timeout"
dask-worker $SCHEDULER_IP:8786 --worker-port=8700:8800 --nanny-port=3000:3100 --nworkers $nprocs --nthreads 1
--death-timeout $death_timeout &
```

entryscript: run.py

https://bitbucket.oci.oraclecorp.com/projects/ODSC/repos/ads-misc/browse/dsc_distributed/run.py?at=distributed

• Dependencies

Base Dependency

```
channels:
- conda-forge
- nodefaults
dependencies:
- conda-forge::dask
- conda-forge::python=3.8
- conda-forge::pip
- pip:
    - oracle-ads
    - supervisor
```

Finding IP

```

from ads.jobs import Job, DataScienceJob, DataScienceJobRun
from ads.common import auth as authutil
import oci
import ads
import ipaddress
import psutil

def what_is_my_ip(jdef_ocid):

    ads.set_auth("resource_principal")

    jobDef = Job.from_datascience_job(jdef_ocid)

    core_client = oci.core.VirtualNetworkClient(**authutil.resource_principal())

    cidr = core_client.get_subnet(jobDef.infrastructure.subnet_id).data.cidr_block
    ipset = [snics[0].address for interface, snics in psutil.net_if_addrs().items()]
    for ip in ipset:
        if ipaddress.ip_address(ip) in ipaddress.ip_network(cidr):
            return ip
    return None

```

User can build images using -

```
ads opctl build-image --backend distributed-job --type dask --image-name <image name> --source <folder> --entrypoint <starting script>
```

This will build the image locally in the users' workstation. The user can then push the image using `docker push` command to ocir

To add dependencies, user will create conda yaml style dependency file and can provide it as additional input in `ads opctl build-image` command.

Eg. User who wants to add lightgbm could provide additional dependency by creating a yaml file such as -

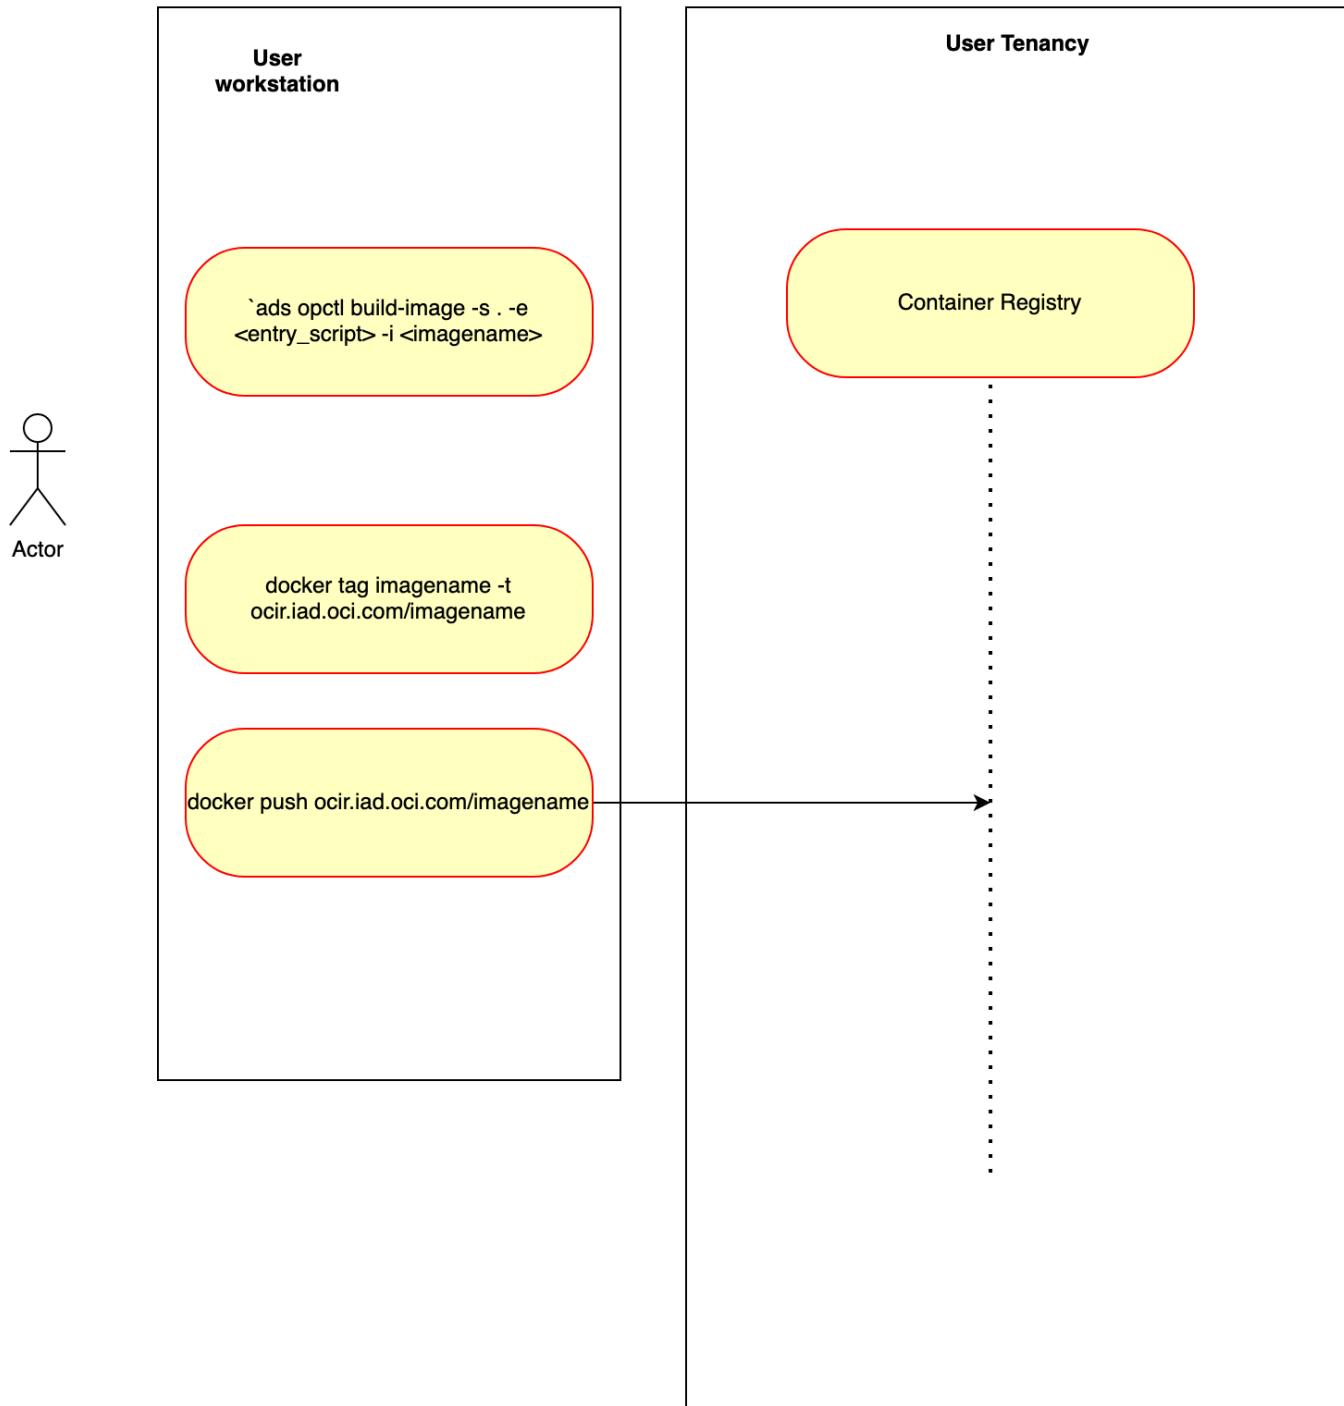
```

channels:
  - conda-forge
  - nodefaults
dependencies:
  - conda-forge::lightgbm

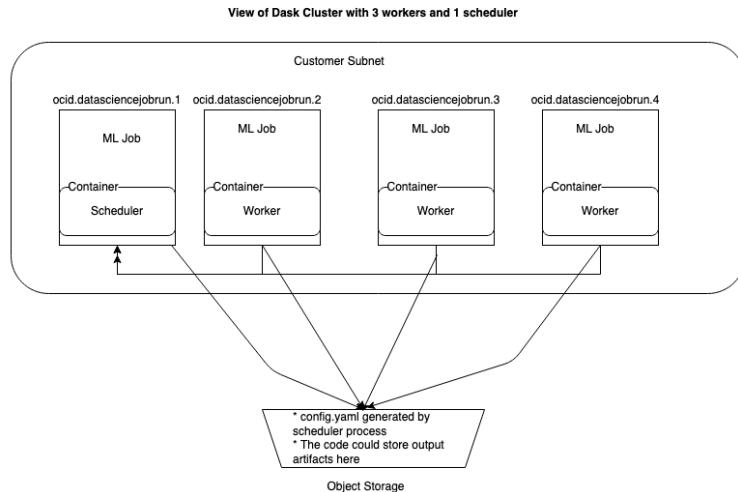
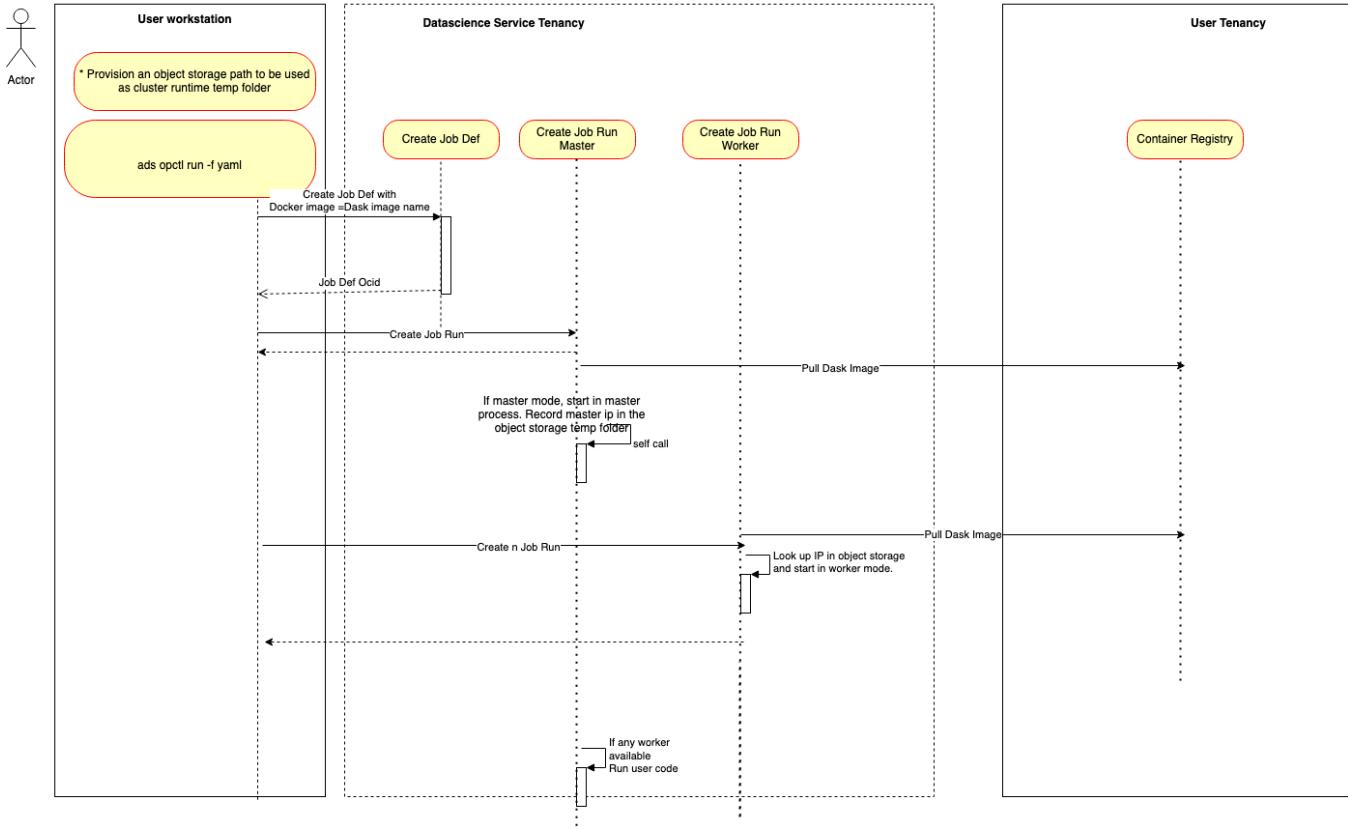
```

The docker image that is built in local is then pushed to ocir registry of the user tenancy

Building and Pushing Docker image for Distributed Training



5.4.1.1.1. Launching Cluster



5.4.1.2. Alternate designs (Designs that were considered but not chosen)

5.4.1.2.1. - With a Driver image

We can build a driver docker image which will launch n jobs with user provided docker command. The driver image will also take ownership of shutting down the cluster based on stopping condition. The ads docker recipe will be restricted to driver code. The user builds and publishes the driver docker image. User will separately build another docker image which contains the code. Such a docker image should reasonably work in the service provided distributed training infrastructure as well. This design is easier to implement in case of horvod than dask. We can probably maintain both approaches. Key advantage of using driver image will be separation of user code and cluster launching logic. When we have service provided distributed training, user can continue to use image corresponding to the code and ignore driver. The driver will not be directly initiated by the user ever. ADS will launch the driver and driver will internally launch the user docker with appropriate configuration.

In case of horvod the worker nodes can automatically discover master node. We could expect user to use horovod public image as their base image.

In case of Dask, the user will have to build the docker image for code from Dockerfile template that we provide. This is because, there is a scheduler node and work node and they are different scripts. So we need some logic to dictate in which mode the image should run.

5.4.1.2.2. Starting workers directly from ADS opctl

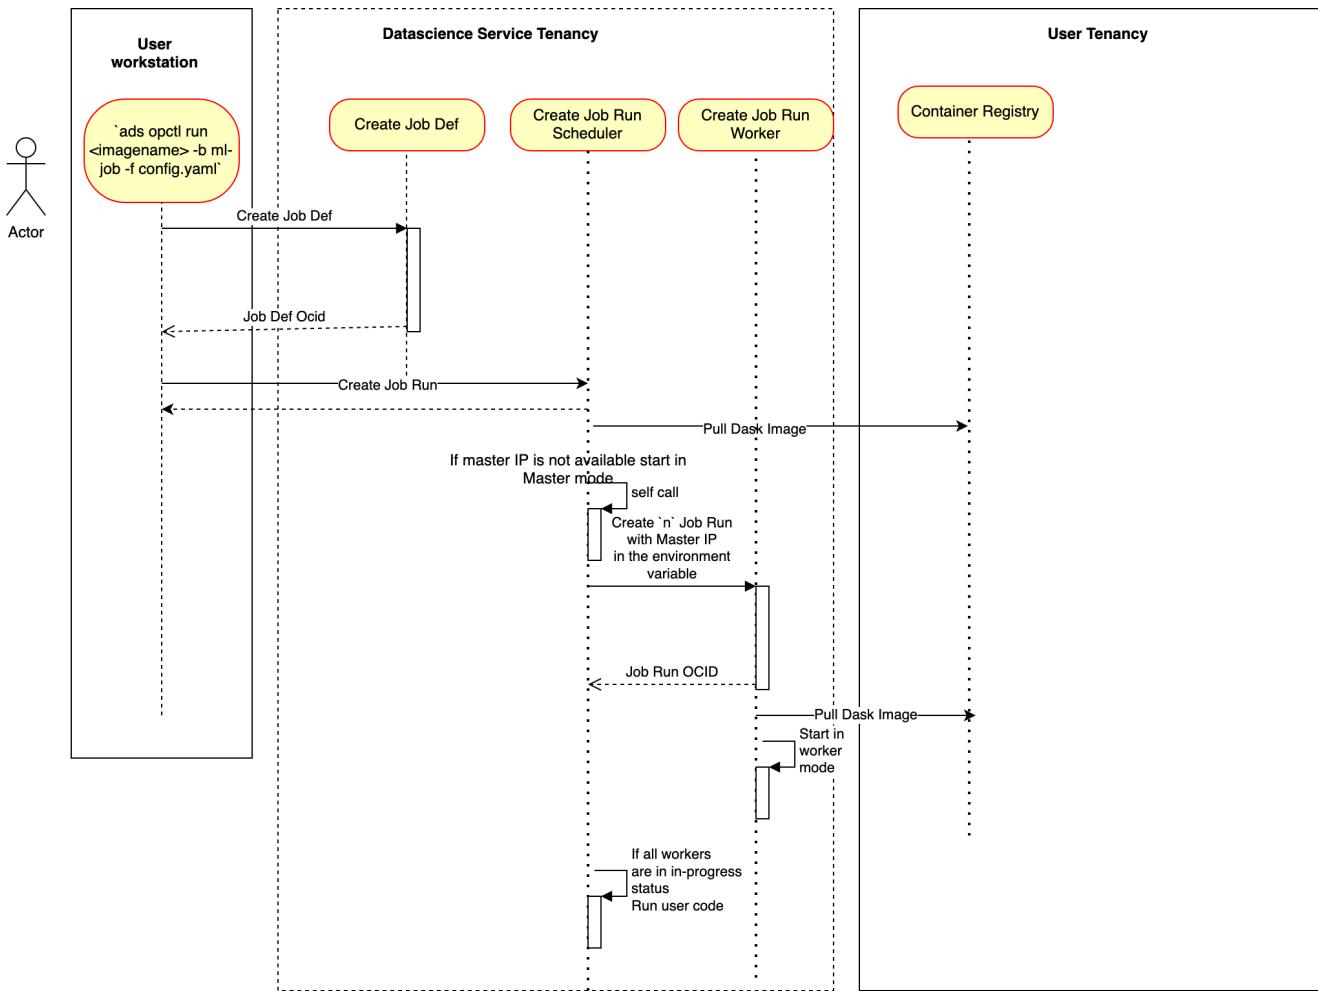
In this approach we start all the workers from the client code (ADS opctl) itself.

5.4.1.2.3. Provide Launch scripts.

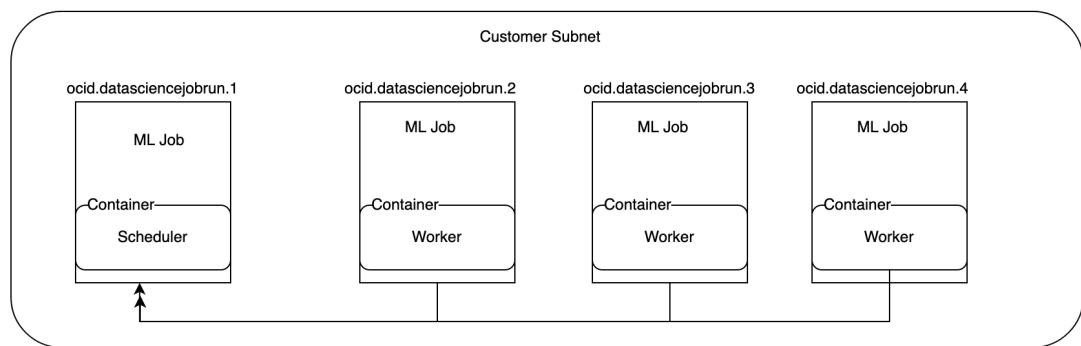
In this approach, we dont tie cluster creation with ADS. We can have launch scripts distributed through our github repo. In this case user will git clone and cd into the distributed-training folder and then follows the Readme.md to build images, push it and then use the launch scripts to create cluster. We can possibly make it easier to dynamically add more workers to the cluster. Updating github is easier than adding features to ADS. Every change to ADS will result in version which may cause confusion to the SDK users who may not be concerned about the `opctl` command line. Adding hacky code to ADS may lead to cause of concern for enterprise users.

Adding temporary features or capabilities to ADS will give poor UX. Whatever we add to ADS should ideally be supported for long term.

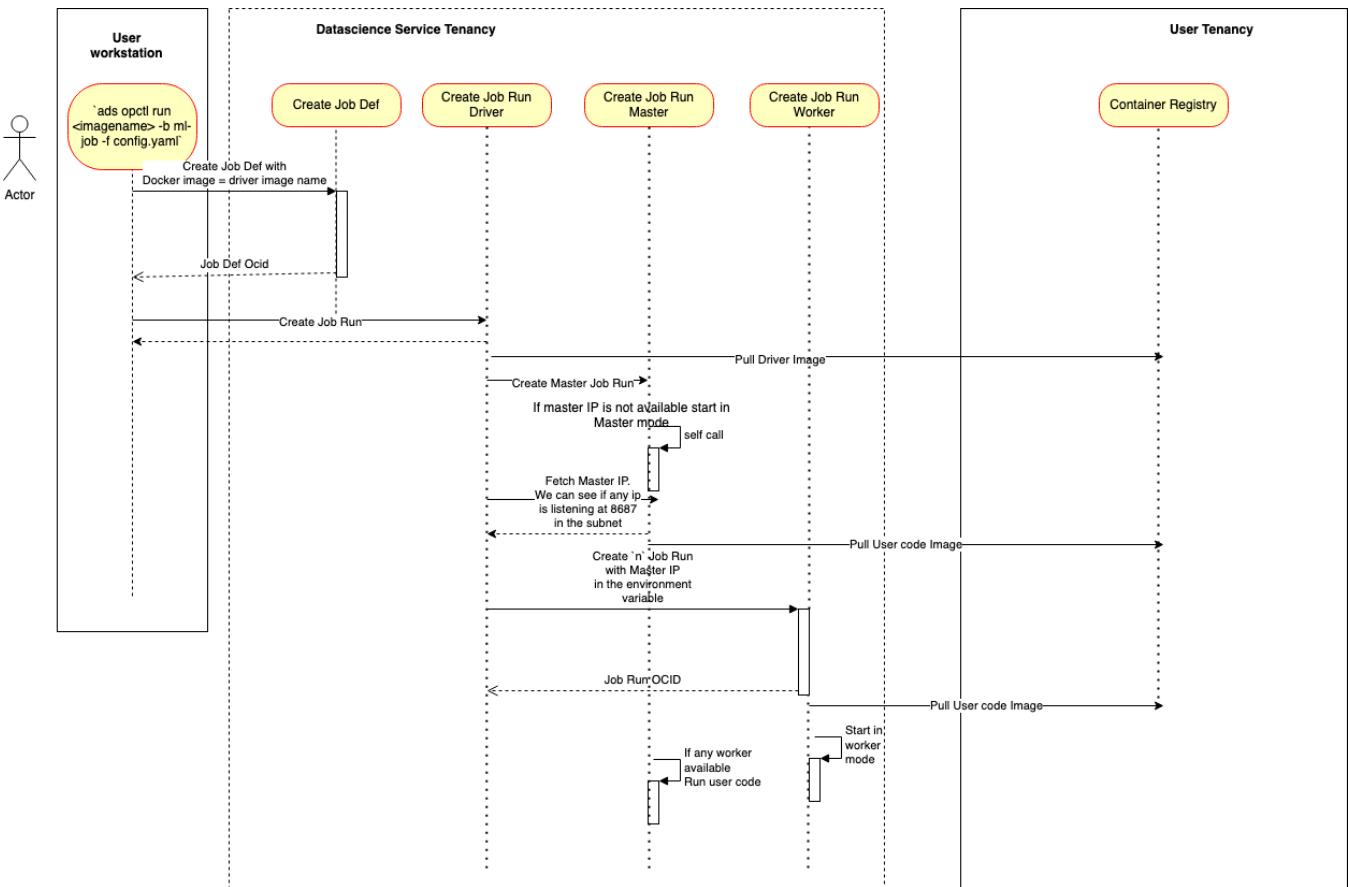
5.4.1.2.4. Scheduler starting n workers.



View of Disk Cluster with 3 workers and 1 scheduler



Launching a cluster with Driver Node



View of Task Cluster with 3 workers and 1 scheduler

